

Sentiment Analysis within Reddit as a Parental Warning Project Milestone

Matthew Nelson

UC Berkeley School of Information
Calgary, Alberta, CA

matthewpeternelson@berkeley.edu

Daniel Wald

UC Berkeley School of Information
Los Angeles, California, USA

daniel.wald@berkeley.edu

Abstract

We examine sentiment analysis on Twitter and Reddit data. This paper explores sentiment analysis methodology used in several papers and the SemEval 2017 challenge in order to provide a parental warning to twitter or subreddit feeds.

Accuracy of individual message analysis is currently 77% (11/14/2017)

1 Credits

Elements of code / models were adapted from the articles cited within the References section, notably (Sentiment Twitter, 2011). This is intended to be a starting point for the structure of our approach and will be adjusted as we progress.

2 Introduction

The internet can be intimidating, especially for parents attempting to limit their children's exposure to explicit content. While the Children's Online Privacy Protection Act (COPA) of 1998 protects children under 13 from exposure to potentially explicit material via login credentials and privacy policy, no viable solution currently exists to further curtail internet access based on dynamic content scoring. To remedy this problem, we build a Convolutional Neural Network to extract a positivity score for individual tweets, reddit pages and subreddits, thus providing an added layer of parental controls for Reddit browsing.

This document is inspired by the recent work to identify trolls that influence democratic republic elections in recent news cycles and the difficulty involved in doing so. Due to the limited scope of a semester Master's class, we have scaled back the scope of our inquiry to be limited to only a select number of Reddit and Twitter instances, asserting a positive / negative score for the session and comments as a data exploration, drawing no conclu-

sion about whether or not the comments / threads are driven by trolls. A wealth of sentiment analysis work on twitter data already exists and much of this contemporary knowledge has been leveraged in the classification sections of this report.

In this paper, we extract Twitter and Reddit messages / replies and feed these into a classifier model which returns one of either two ("Positive" and "Negative" in the case of the STS Tweet Corpus Training Data (Sentiment Twitter, 2011)) or three classes ("Positive", "Negative" and "Neutral" in the case of the SemEval Challenge Task4a Training Data). We experiment with 2 models and feature engineering to modify weights and biases. The models used include Convolutional Neural Network (CNN), herein referred to as "baseline model" and Recurring Neural Network (RNN). Both models are driven by the text contained within the posts themselves. We extract interaction variables from the message itself in the form of number of interactions and number of chained parents (i.e. a 3rd order comment is in response to a 2nd order comment which is in turn in response to a 1st order (top level) comment or posting).

This paper is organized as follows. Section 3 outlines the methodology with subsections detailing selection of corpus and models. Section 4 outlines results of model with subsections detailing modelling accuracy as well as accuracy when compared against Mechanical Turk assessment. Section 5 will discuss conclusions and path forward for future enhancements. Section 6 includes references.

3 Methodology

Data is scraped from the Twitter feed, reddit submission or subreddit and fed into the model, described in subsections, below. The first 100 *CONFIRM* comments and replies (limited to reduce page load latency) contents of the page are then

fed into the model on a dynamic *CONFIRM* basis which creates a score. This score is then displayed directly on the page being called via a Google Chrome Browser Extension. *CONFIRM*. Posts and pages scoring $< 30\%$ positive are then marked as toxic / extremely negative using a simple data visualization or an emoji *CONFIRM*.

The model was then scored for accuracy using mechanical turks to assess positive / negative nature of each page with a 60% *CONFIRM accuracy rate measured against the post and a 85% CONFIRM accuracy rate measured when determining whether or not the page over-all communicated a positive or negative message*. Possible errors include too small a training corpus, varying degrees of decorum among turks and the model itself as the models explored in this paper lack of capability to discern complex communication such as sardonic wit, sarcasm or meme driven comments.

3.1 Corpus Selection

The use of deep learning neural networks within the keras library (Using Tensorflow at the backend) allows us to more easily fit and optimize the model rather than struggle with matrix multiplication.

Our baseline model was initially trialed on the following corpuses: *need table reference - need to update table as well*

Corpus	Entries	BL Score
SemEval 2017	10,000	52%
SemEval 2012-2017	23,000	45%
STS Tweet Corpus	1,600,000	77%

The STS corpus is a collection of 1.6 million English tweets with positive and negative sentiment labels. These labels were generated by requesting queries through the Twitter search API for tweets with positive and negative emoticons in their content. The emoticons are removed from the content.

These corpuses score microblog social media posts as either a two class 'positive' / 'negative' or a three class 'positive' / 'negative' / 'neutral'.
ADD REFERENCES

3.2 Baseline Model: CNN

Convolutional Neural Network was used as a baseline model in an attempt to replicate results of SemEval 2017 papers (Sentiment Twitter, 2011).

The model was trained using the SemEval 2017 corpus of 10,000 manually categorized microblog

posts. Initial execution of returned approximately 52% accuracy within test data and, according to this reader, a one-sided prediction when applied to Twitter and Reddit posts. We joined our corpus with the last 5 years of the SemEval challenge corpuses in order to grow our corpus into a size (23,000 rows) that support the use of neural networks, yielding even worse performance. Based on poor response from limited data and an understanding that neural networks require hundreds of thousands to millions of entries, we decided to leverage a larger social media generated data set, the STS Tweet corpus. This corpus contains 1.6 Million Tweets which were auto-labeled based on pre-defined 'positive' or 'negative' emoticons present in the original tweet.

3.3 RNN Model

We believe the chained nature of comments lends itself to a recurring neural network and that incorporation of this model will result in higher accuracy than with the baseline model. Work on RNN has not yet started.

I expect this section to take up a full half page

3.4 Feature Engineering: Weights and Biases

Multiple levels of feature engineering were used on the raw tweets in order to ready them for ingestion into the CNN model.

Cleaning: Only the Sentiment and Tweet Text columns from the corpora were utilized, the remainder were removed including 'ItemID', 'Date-Time', 'Query', and 'SentimentSource'. **More to come.**

Tokenization: Tweets were parsed into individual tokens (words) and forced to lowercase. **More to come.**

Word Vectorization: Google's Word2Vec model was utilized to generate a vector representation of each word, using 100 dimensional vectors and setting a minimum threshold on the word count for eligible words in order to prevent rare words from entering the vector.

TF-IDF Word Score: ScikitLearn's TfidfVectorizer was used to generate a TF-IDF representation of the vocabulary of the STS corpus to provide a more accurate representation of each word's importance across the corpora.

Tweet Vector: An overall vector representation of each Word was created by multiplying the individual Word vector by the same Word's TF-IDF




Figure 1: STS corpus where 0 = negative and 4 = positive. Note, this plot is a place holder to capture sample latex code - we will replace with better data visualizations for final

score. This multiplied word vector was then incorporated into an average Tweet vector that would be used as our input feature to the Convolutional Neural Network Sentiment Classifier Model.

4 Results

Discussion of final results

4.1 Displaying the Data

Data Visualization discussion / pictures. How do we meet the expectations in the introduction?

Only plot done so far is showing the count of each sentiment classification in the raw STS corpus. Don't know how to upload that into the document though...

Figure 1, above, shows the STS corpus distribution between positive and negative tweets. As you can see, the distribution is even with approximately 800k positive and 800k negative tweets.

4.2 Modelling Accuracy

The accuracy of the initial CNN model on the STS corpus and utilizing Tweet Vectors as the input feature provides an accuracy of 77%. This has not yet been optimized and is very much a baseline.

4.3 Human Accuracy

Assess snapshot of predictions using mechanical turks to obtain unbiased approach for individual posts as well as full pages / thread level positivity / negativity. Note there will be some jitter in these results due to randomized differences in turks as well as a self selection bias that may reflect the social proclivities of those that participate in the mechanical turk marketplace.

5 Conclusion

No conclusions yet.

5.1 Future Work

Integrate the predicted Sentiment (output of CNN) for each Reddit comment.

Create a separate model (RNN) which attempts to predict the sentiment of the next child comment from the sentiment of the previous parent comments in a single post on a Reddit page.

Perhaps limit project focus to Reddit Pages and drop focus on Twitter User History?

What would we do if we had more time? This is where we can address shortcomings or budgetary issues

Acknowledgments

The acknowledgments should go immediately before the references. Do not number the acknowledgments section. Do not include this section when submitting your paper for review.

References

- Socher, Perelygin, et al. 2013. *Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank*, Conference on Empirical Methods in Natural Language Processing. <http://acl2015.org/publication.html>
- Agarwal, Xie, Vovsha, Rambow, Passonneau. 2011. *Sentiment analysis of Twitter data*. LSM, Portland, OR. <https://dl.acm.org/citation.cfm?id=2021114>
- Cliche. 2017. *BB twtr at SemEval-2017 Task 4: Twitter Sentiment Analysis with CNNs and LSTMs*. SemEval, 2017 <https://aclanthology.info/pdf/S/S17/S17-2094.pdf>
- Balicas. 2017. *TwisE at SemEval-2017 Task 4: Five-point Twitter Sentiment Classification and Quantification*. SemEval, 2017 <https://aclanthology.info/pdf/S/S17/S17-2127.pdf>